

# Learning to Optimize in Swarms

Yue Cao<sup>1</sup>, Tianlong Chen<sup>2</sup>, Zhangyang Wang<sup>2</sup> and Yang Shen<sup>1</sup>

<sup>1</sup>Department of Electrical and Computer Engineering, and <sup>2</sup>Department of Computer Science and Engineering, Texas A&M University, College Station, TX 77843, United States.

## Abstract

- **Learning to optimize** has emerged as a powerful framework for various optimization tasks.
- Current such “meta-optimizers” often learn in the space of continuous optimization algorithms that are **point-based** and **uncertainty-unaware**.
- We learn in an extended space of **both** point-based and **population-based** optimization algorithms.
- We incorporate a Boltzmann-shaped **posterior over the global optimum** into meta-loss to balance the exploitation-exploration trade-off during search.
- Empirical results over non-convex test functions and the **protein docking** application demonstrate that this new meta-optimizer outperforms existing competitors.

## Methods

- **Updating Rules:** Iterative optimization algorithms, either point-based or population-based, share a generic expression of update formulae:

$$\mathbf{x}^{t+1} = \mathbf{x}^t + \delta \mathbf{x}^t$$

The update is often a function  $g(\cdot)$  of the historic sample values, objective values, and gradients. For instance, in particle swarm optimization (PSO), we have

$$\begin{aligned} \delta \mathbf{x}_j^t &= g(\{\mathbf{x}_j^\tau, f(\mathbf{x}_j^\tau), \nabla f(\mathbf{x}_j^\tau)\}_{j=1, \tau=1}^{k, t}) \\ &= w \delta \mathbf{x}_j^{t-1} + r_1(\mathbf{x}_j^t - \mathbf{x}_j^{t*}) + r_2(\mathbf{x}_j^t - \mathbf{x}^{t*}) \end{aligned}$$

In **our** approach, we parameterize the update rule  $g(\cdot)$  through RNN, and introduce **intra-** and **inter-particle attention mechanisms**:

$$g_i(\cdot) = \text{RNN}_i(\alpha_i^{\text{inter}}(\{\alpha_j^{\text{intra}}(\{\mathbf{S}_j^\tau\}_{\tau=1}^t)\}_{j=1}^k), \mathbf{h}_i^{t-1})$$

- **Population-based and Point-based Features:** Inspired from both point- and population-based algorithms, we choose the following four features for particle  $i$  at iteration  $t$ :
  - gradient:  $\nabla f(\mathbf{x}_i^t)$
  - momentum:  $\mathbf{m}_i^t = v_{\tau=1}^t(1 - \beta)\beta^{t-1}\nabla f(\mathbf{x}_i^t)$
  - velocity:  $\mathbf{v}_i^t = \mathbf{x}_i^t - \mathbf{x}_i^{t*}$
  - attraction:  $\frac{\sum_j (e^{-\alpha d_{ij}}(\mathbf{x}_i^t - \mathbf{x}_j^t))}{\sum_j e^{-\alpha d_{ij}}}$ , for all  $j$  that  $f(\mathbf{x}_j^t) < f(\mathbf{x}_i^t)$ .  $\alpha$  is the hyperparameter and  $d_{ij} = \|\mathbf{x}_i^t - \mathbf{x}_j^t\|_2$ .

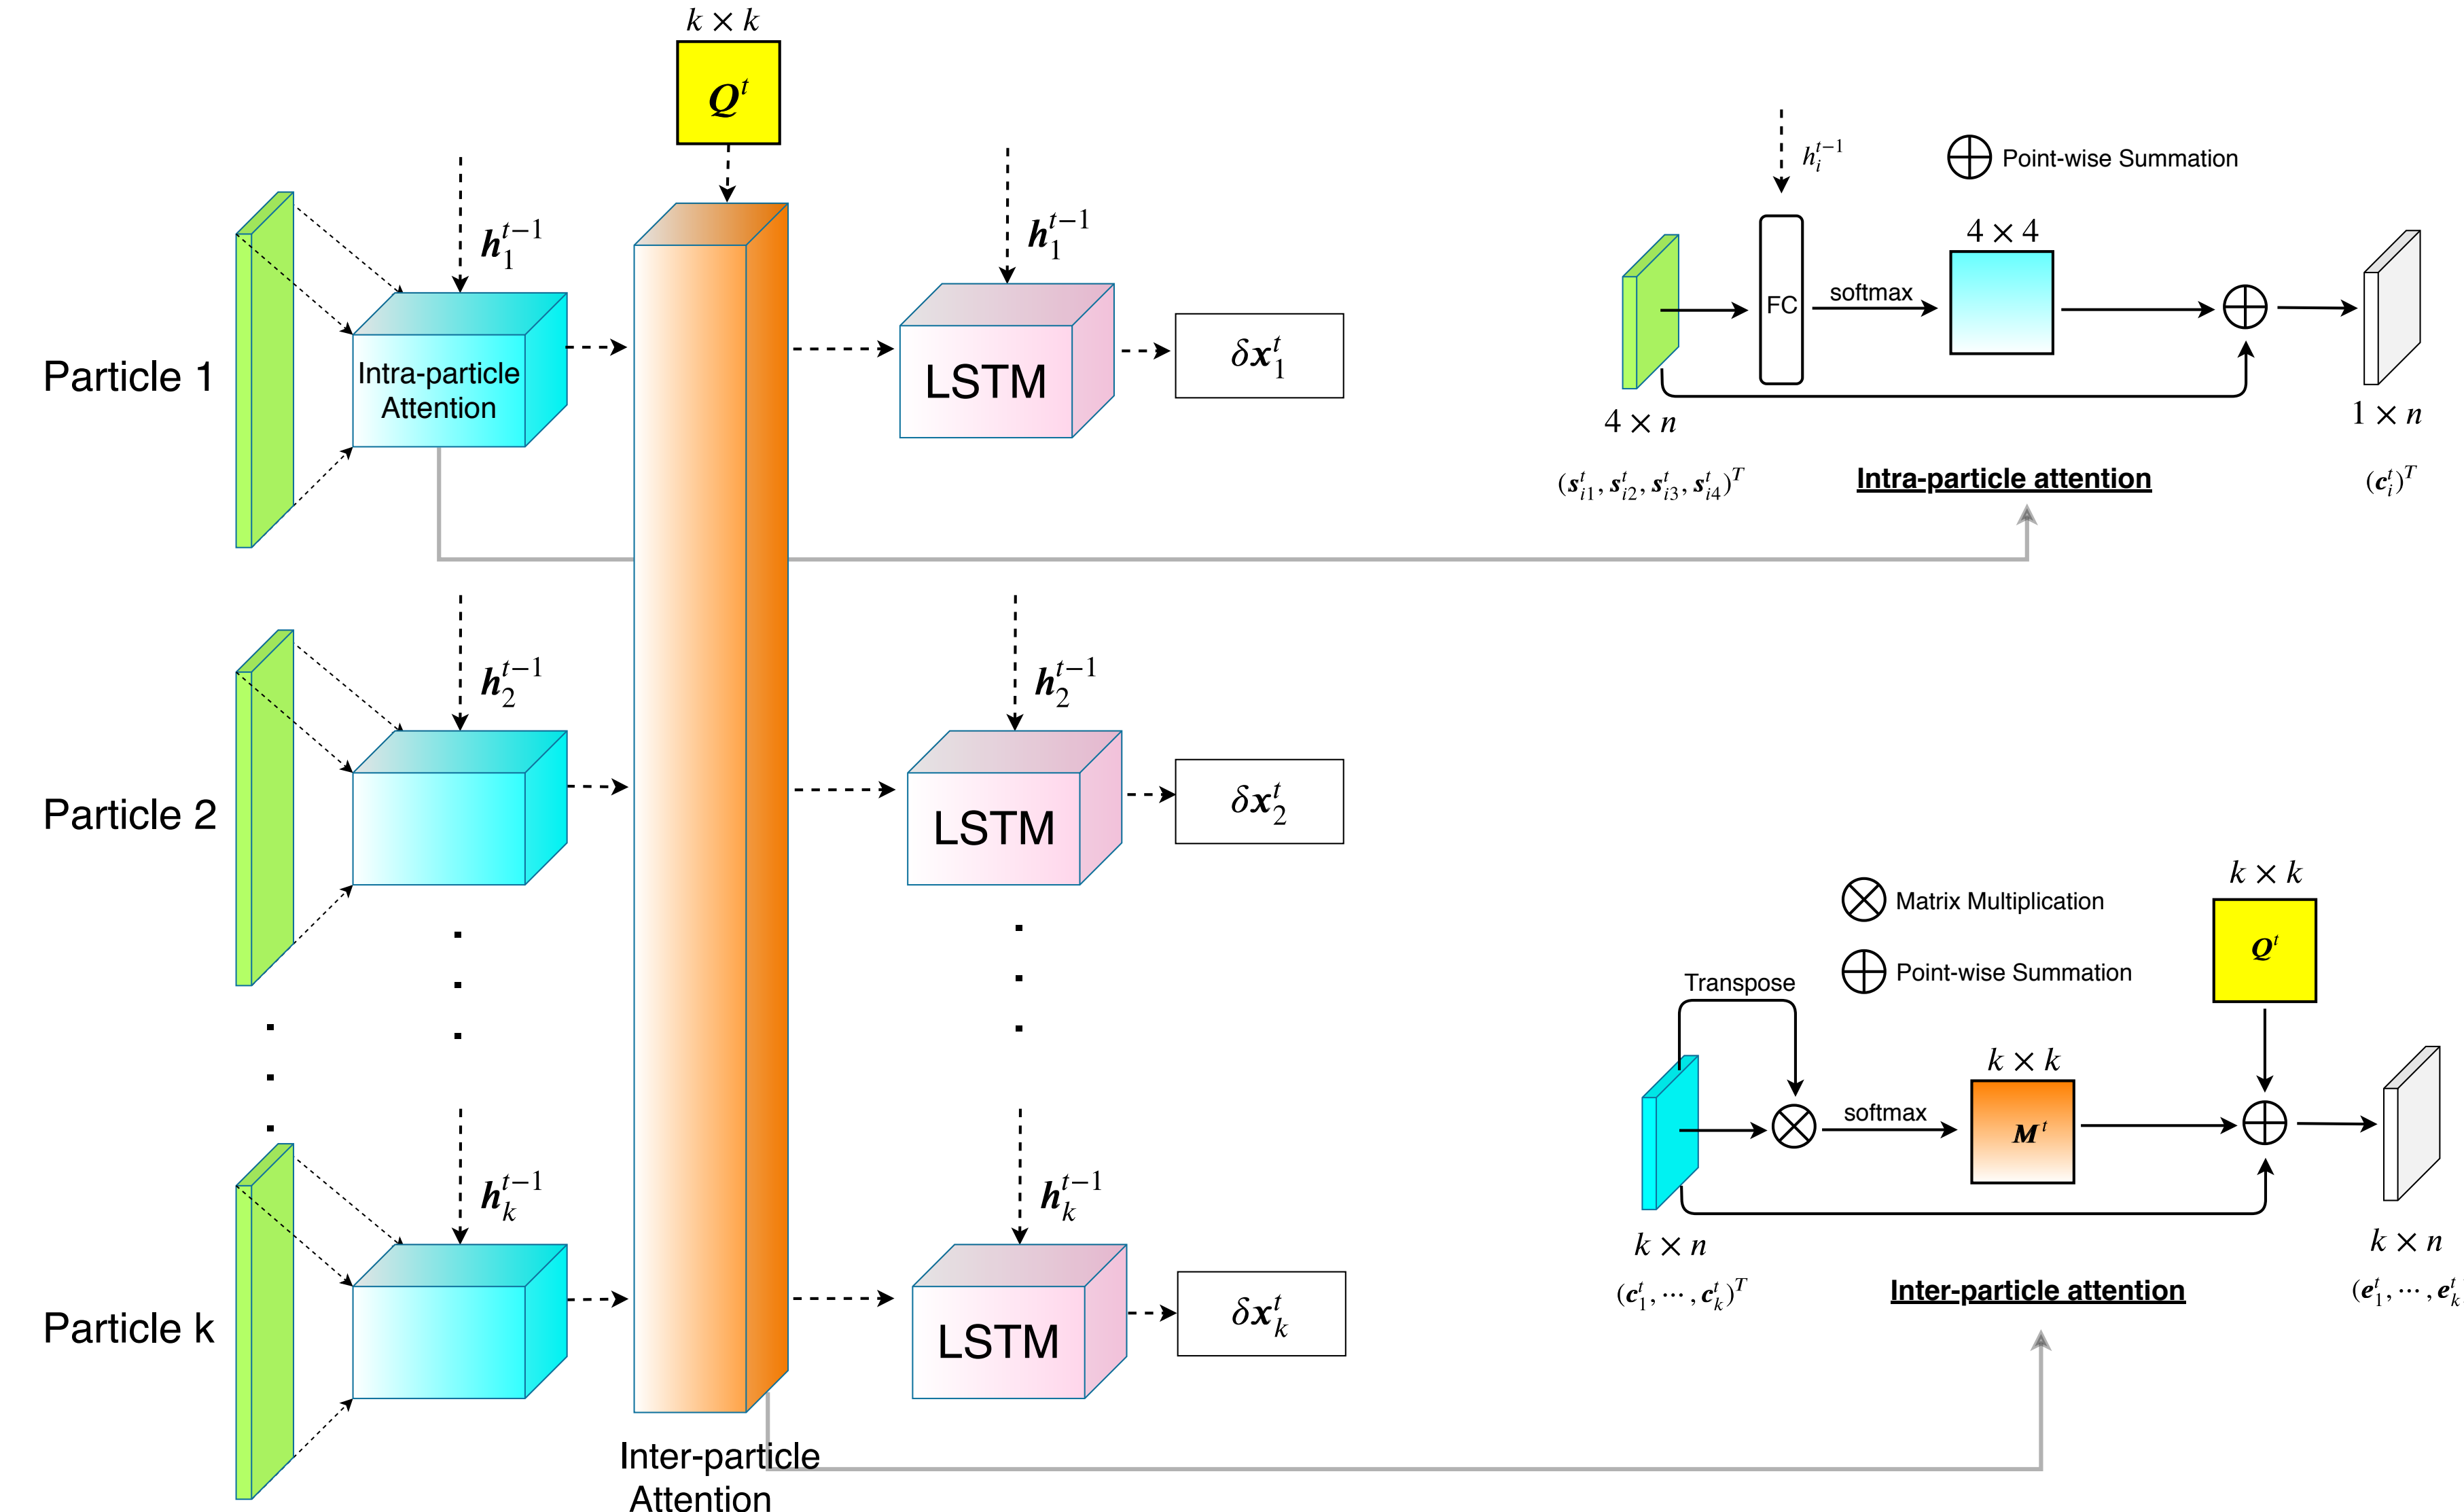
- **Loss Function:** In order to balance the exploration-exploitation tradeoff, we combine the cumulative regret and the entropy of the posterior over the global optimum:

$$\ell_f(\phi) = \frac{T}{t=1} \sum_{j=1}^k f(\mathbf{x}_j^t) + \lambda h(p(\mathbf{x}^* |_{t=1}^T D_t)),$$

where the posterior is a **Boltzmann distribution** [3]:

$$p(\mathbf{x}^* |_{t=1}^T D_t) \propto \exp(-\rho \hat{f}(\mathbf{x}))$$

## Overall architectures and attention modules




- **Intra-particle (feature-level) attention:**  $b_{ij}^t = \mathbf{v}_a^T \tanh(\mathbf{W}_a \mathbf{s}_{ij}^t + \mathbf{U}_a \mathbf{h}_{ij}^t)$ ,  $p_{ij}^t = \frac{\exp(b_{ij}^t)}{\sum_{r=1}^4 \exp(b_{ir}^t)}$ ,  $\mathbf{c}_i^t = \sum_{r=1}^4 p_{ir}^t \mathbf{s}_{ir}^t$
- **Inter-particle (sample-level) attention:**  $\mathbf{e}_j^t = \gamma \sum_{r=1}^k m_{rj}^t q_{rj}^t \mathbf{c}_r^t + \mathbf{c}_j^t$

## Test Function Results


**LOIS** outperforms DM\_LSTM [1] and hand-engineered algorithms in optimizing non-convex Rastrigin functions:

$$f(\mathbf{x}) = \sum_{i=1}^n x_i^2 - \frac{n}{i=1} \alpha \cos(2\pi x_i) + \alpha n$$



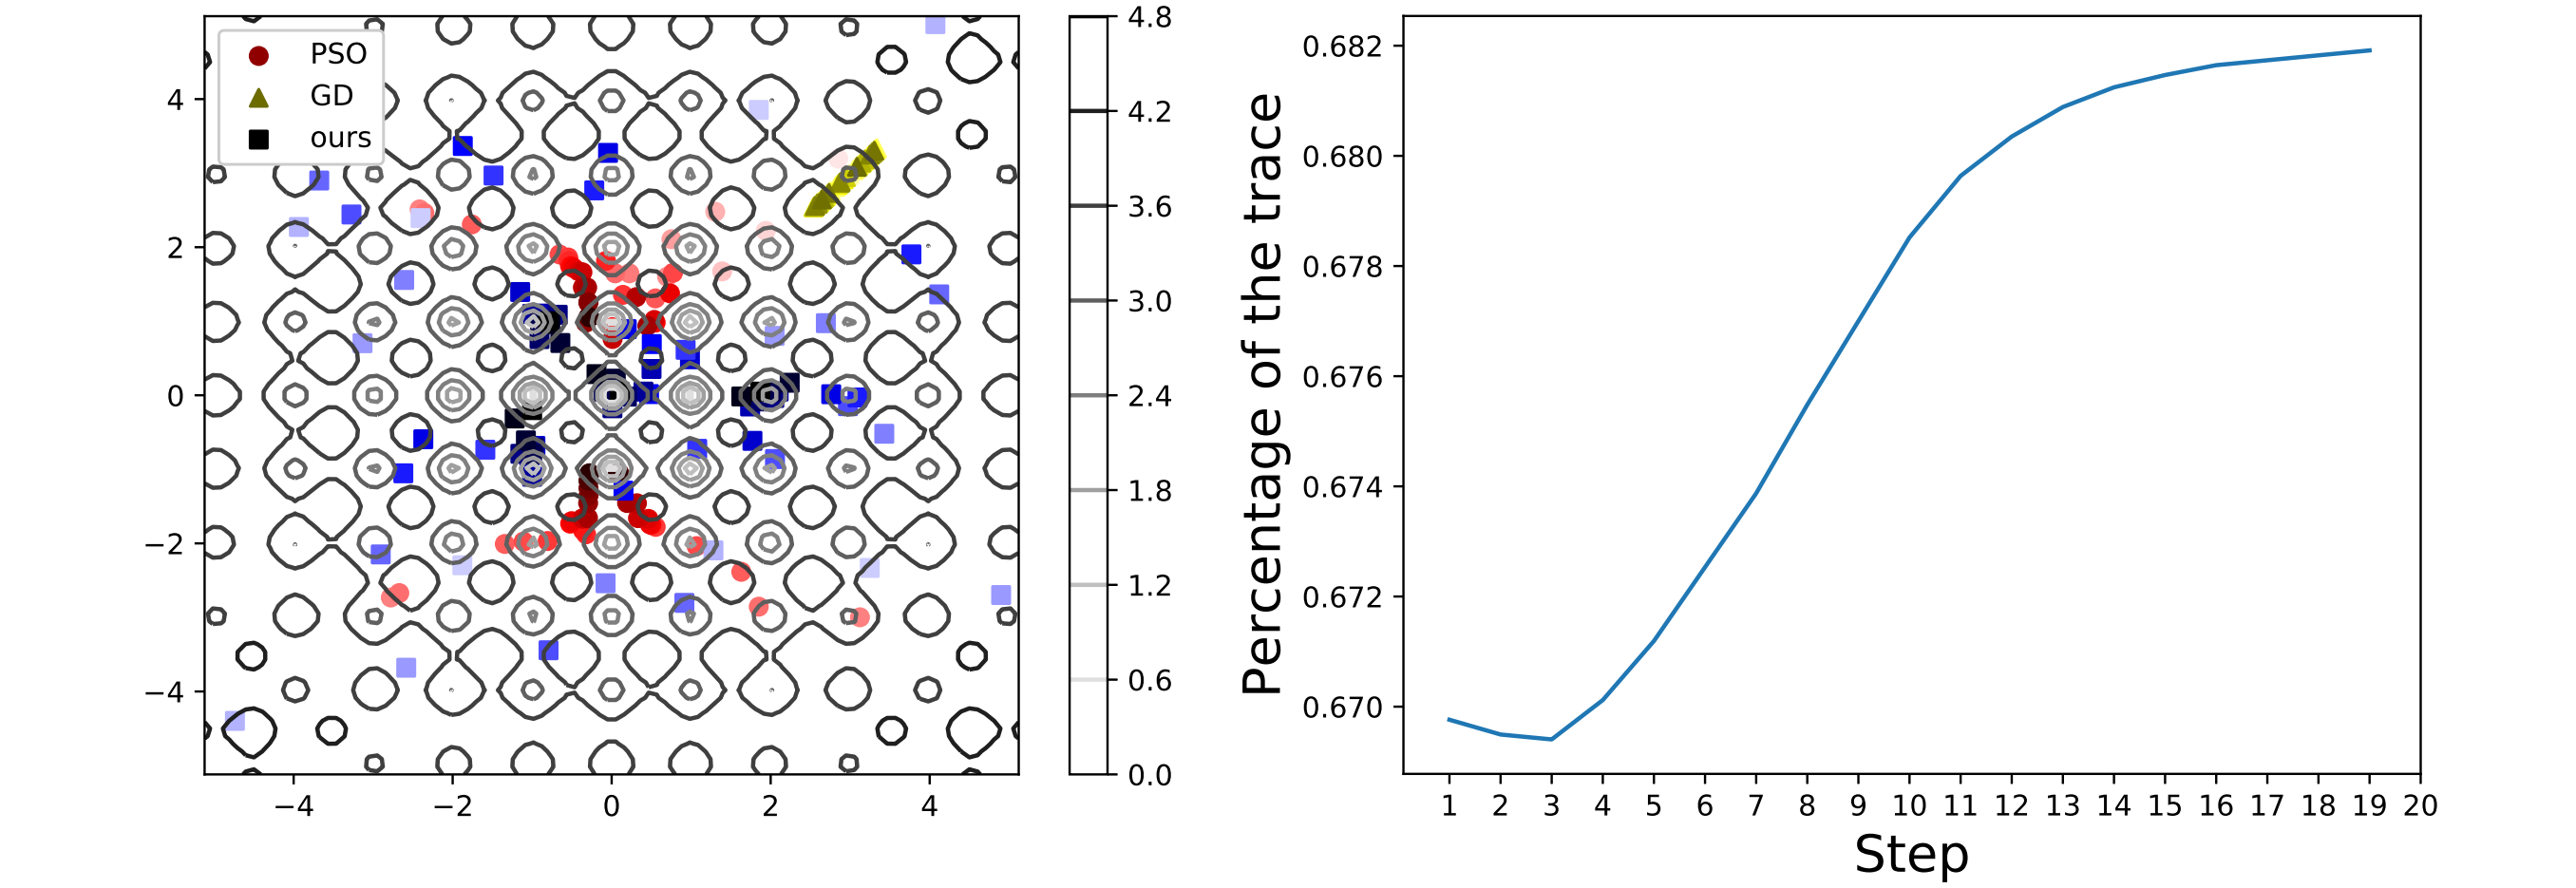
## Protein Docking Results

*Ab initio* protein docking represents a major challenge for optimizing a noisy and costly function in a high-dimensional space [3]. We parameterize the search space as  $\mathbb{R}^{12}$  as in [3]. The corresponding  $f(\mathbf{x})$  (energy function) is fully differentiable. **LOIS** outperforms PSO in energy scores for three protein-protein pairs of different difficulty-levels.



## Interpretation Results

- The trace only accounts for 66%-69% over iterations as shown in (b). This demonstrates the importance of collaboration, a unique advantage of population-based algorithms.



(a) Paths of the first 80 samples of our meta-optimizer, PSO and GD for the 2D Rastrigin function. (b) The percentage of the trace of  $\gamma \mathbf{Q}^t \odot \mathbf{M}^t + \mathbf{I}$  (reflecting self-impact on updates) over iteration  $t$ .

- In the first 6 iterations, the population-based features (3 & 4) contribute to the update the most. Point-based features (1 & 2) start to play an important role later:




Figure: Feature distribution over the first 20 iterations for our meta-optimizer.

## Ablation Study

Dimension	<b>B<sub>0</sub></b>	<b>B<sub>1</sub></b>	<b>B<sub>2</sub></b>	<b>B<sub>3</sub></b>	<b>Proposed</b>
10	55.4±13.5	48.4±10.5	40.1±9.4	20.4±6.6	12.3±5.4
20	140.4±10.2	137.4±12.7	108.4±13.4	48.5±7.1	43.0 ±9.2

Table: **B<sub>0</sub>**: the DM\_LSTM baseline. **B<sub>1</sub>**: running DM\_LSTM for  $k$  times and choosing the best solution. **B<sub>2</sub>**: using  $k$  independent particles, each with the two point-based features and the intra-particle attention module. **B<sub>3</sub>**: adding the two population-based features and the inter-particle attention module to **B<sub>2</sub>**. **Proposed**: adding an entropy term in meta loss to **B<sub>3</sub>**.

## Acknowledgement

This work is in part supported by the National Institutes of Health (R35GM124952 to YS). Part of the computing time is provided by the Texas A&M High Performance Research Computing.

## References

- [1] Marcin Andrychowicz et al. (2016). 'Learning to learn by gradient descent by gradient descent.' In Advances in Neural Information Processing Systems, pages 3981–3989.
- [2] Yue Cao, Tianlong Chen, Zhangyang Wang, Yang Shen. (2019). 'Learning to Optimize in Swarms'. NeurIPS 2019
- [3] Yue Cao and Yang Shen. (2019). 'Bayesian active learning for optimization and uncertainty quantification in protein docking.' arXiv preprint arXiv:1902.00067